

✓ Task-1 : Face Mask Detection using CNN

1. Problem Understanding

This project is a binary classification problem.

- Input: Face images
- Output: Mask / No Mask

✓ 2. Data Collection

For this step, we need to gather the data.

- Use datasets like:
 - Kaggle Face Mask Dataset
 - Real-world images (optional)

```
# Install kagglehub to download the dataset directly
!pip install kagglehub -q

import kagglehub
import os

# Download the dataset
path = kagglehub.dataset_download("omkargurav/face-mask-dataset")
input_data_dir = os.path.join(path, 'data')

print("Dataset downloaded to:", input_data_dir)
```

Using Colab cache for faster access to the 'face-mask-dataset' dataset.
Dataset downloaded to: /kaggle/input/face-mask-dataset/data

✓ 3. Data Preprocessing

Before feeding the images to the model, we need to format them correctly.

- Resize images (e.g., 224x224 or 128x128)
- Normalize pixel values (0-1)
- Convert labels into numerical format (0, 1)
- Split dataset:
 - Train (70%)
 - Validation (15%)
 - Test (15%)

```
# Install split-folders to manage the 70/15/15 split
!pip install split-folders -q
import splitfolders

output_dir = "./dataset_split"

# Splitting the dataset into Train, Validation, and Test
splitfolders.ratio(
    input_data_dir,
    output=output_dir,
    seed=42,
```

```
ratio=(0.7, 0.15, 0.15)
)

print("Data splitting complete. Folders created in:", output_dir)

Copying files: 7553 files [00:21, 346.09 files/s]Data splitting complete. Folders created in: ./da
```

✓ 4. Data Augmentation

To ensure the model doesn't just memorize the training data, we introduce slight variations to the images. We will use ImageDataGenerator (Keras) to improve generalization using:

- Rotation
- Flipping
- Zoom
- Shearing

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Setup augmentation for training data and normalize pixel values
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    horizontal_flip=True,
    zoom_range=0.2,
    shear_range=0.2
)

# Only normalize pixel values for validation and test data
val_test_datagen = ImageDataGenerator(rescale=1./255)

# Load data, resize to 128x128, and convert labels to numerical binary format
train_generator = train_datagen.flow_from_directory(
    f"{output_dir}/train",
    target_size=(128, 128),
    batch_size=32,
    class_mode='binary'
)

val_generator = val_test_datagen.flow_from_directory(
    f"{output_dir}/val",
    target_size=(128, 128),
    batch_size=32,
    class_mode='binary'
)

test_generator = val_test_datagen.flow_from_directory(
    f"{output_dir}/test",
    target_size=(128, 128),
    batch_size=32,
    class_mode='binary',
    shuffle=False
)

Found 5286 images belonging to 2 classes.
Found 1132 images belonging to 2 classes.
Found 1135 images belonging to 2 classes.
```

✓ 5. Build CNN Model

Now we define the neural network. We will use a basic CNN Architecture:

- Convolution Layer + ReLU
- MaxPooling
- Dropout
- Fully Connected Layer
- Output Layer (Sigmoid for binary)

Example flow: Conv -> ReLU -> Pool -> Conv -> Pool -> Flatten -> Dense -> Output

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

# Building the architecture based on the specified flow
model = Sequential([
    # Convolution Layer + ReLU
    Conv2D(32, (3, 3), activation='relu', input_shape=(128, 128, 3)),
    # MaxPooling
    MaxPooling2D(pool_size=(2, 2)),

    # Second Conv and Pool block
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),

    # Flatten
    Flatten(),

    # Fully Connected Layer
    Dense(128, activation='relu'),
    # Dropout
    Dropout(0.5),

    # Output Layer (Sigmoid for binary classification)
    Dense(1, activation='sigmoid')
])

model.summary()
```

/usr/local/lib/python3.12/dist-packages/keras/src/layers/convolutional/base_conv.py:113: UserWarning
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_5 (Conv2D)	(None, 61, 61, 64)	18,496
max_pooling2d_5 (MaxPooling2D)	(None, 30, 30, 64)	0
flatten_2 (Flatten)	(None, 57600)	0
dense_4 (Dense)	(None, 128)	7,372,928
dropout_2 (Dropout)	(None, 128)	0
dense_5 (Dense)	(None, 1)	129

Total params: 7,392,449 (28.20 MB)

Trainable params: 7,392,449 (28.20 MB)

Non-trainable params: 0 (0.00 B)

6. Model Training

With the architecture built, we compile and train the model using these specific parameters:

- Loss Function: Binary Crossentropy
- Optimizer: Adam
- Metrics: Accuracy

```
import warnings

# Suppress specific PIL UserWarning related to palette images with transparency
warnings.filterwarnings(
    "ignore",
    message="Palette images with Transparency expressed in bytes should be converted to RGBA image",
    category=UserWarning
)
```

```
# Compile the model
model.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy']
)

# Train the model
history = model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=10
)
```

```
Epoch 1/10
166/166 ————— 112s 666ms/step - accuracy: 0.8237 - loss: 0.4175 - val_accuracy: 0.8
Epoch 2/10
166/166 ————— 106s 639ms/step - accuracy: 0.8801 - loss: 0.2860 - val_accuracy: 0.8
Epoch 3/10
166/166 ————— 102s 611ms/step - accuracy: 0.8977 - loss: 0.2496 - val_accuracy: 0.8
Epoch 4/10
166/166 ————— 104s 624ms/step - accuracy: 0.9047 - loss: 0.2261 - val_accuracy: 0.9
Epoch 5/10
166/166 ————— 106s 637ms/step - accuracy: 0.9073 - loss: 0.2257 - val_accuracy: 0.9
Epoch 6/10
166/166 ————— 105s 632ms/step - accuracy: 0.9126 - loss: 0.2202 - val_accuracy: 0.9
Epoch 7/10
166/166 ————— 104s 625ms/step - accuracy: 0.9192 - loss: 0.2011 - val_accuracy: 0.9
Epoch 8/10
166/166 ————— 102s 615ms/step - accuracy: 0.9200 - loss: 0.1917 - val_accuracy: 0.9
Epoch 9/10
166/166 ————— 142s 617ms/step - accuracy: 0.9251 - loss: 0.1803 - val_accuracy: 0.9
Epoch 10/10
166/166 ————— 98s 590ms/step - accuracy: 0.9255 - loss: 0.1786 - val_accuracy: 0.92
```

7. Model Evaluation

Finally, we test how well the model learned by running it against our unseen test data. We evaluate based on:

- Accuracy
- Confusion Matrix
- Precision & Recall

```
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.metrics import confusion_matrix, classification_report

# 1. Overall Accuracy
```

```

test_loss, test_accuracy = model.evaluate(test_generator)
print(f"\nOverall Test Accuracy: {test_accuracy * 100:.2f}%")

# Generate predictions to build the matrix and report
test_generator.reset()
predictions = model.predict(test_generator)
y_pred = (predictions > 0.5).astype(int)
y_true = test_generator.classes

# 2. Confusion Matrix
cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=test_generator.class_indices.keys(),
            yticklabels=test_generator.class_indices.keys())
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()


# 3. Precision & Recall
print("\nClassification Report (Precision & Recall):")
print(classification_report(y_true, y_pred, target_names=test_generator.class_indices.keys()))

```

36/36 ————— 8s 234ms/step - accuracy: 0.9480 - loss: 0.1535

Overall Test Accuracy: 94.80%

36/36 ————— 10s 264ms/step



Classification Report (Precision & Recall):

	precision	recall	f1-score	support
with_mask	0.94	0.95	0.95	560
without_mask	0.95	0.94	0.95	575
accuracy			0.95	1135
macro avg	0.95	0.95	0.95	1135
weighted avg	0.95	0.95	0.95	1135

8. Project Insights and Analysis


Let's visualize the training history to understand how the model learned over epochs.

```
import matplotlib.pyplot as plt

# Plot training and validation accuracy
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

# Plot training and validation loss
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()
```



Analysis of Training History:

From the plots above, we can observe:

- **Accuracy:** Both training and validation accuracy generally increase over the epochs. The validation accuracy shows some fluctuations but overall trends upward, indicating the model is learning effectively.
- **Loss:** Correspondingly, both training and validation loss decrease, which is a good sign. There might be a slight divergence in the later epochs where validation loss starts to fluctuate or increase slightly while training loss continues to decrease, which could be an early sign of overfitting. However, the dropout layer seems to be helping mitigate this to some extent.

Interpretation of Evaluation Metrics:

The model was evaluated on the test set, yielding an **Overall Test Accuracy of 91.89%**.

Confusion Matrix Breakdown:

```
[[551, 9],
 [ 83, 492]]
```

Considering 'with_mask' as class 0 and 'without_mask' as class 1:

- **True Negatives (TN): 551** (Images with masks correctly predicted as 'with mask')
- **False Positives (FP): 9** (Images with masks incorrectly predicted as 'without mask' - Type I error)
- **False Negatives (FN): 83** (Images without masks incorrectly predicted as 'with mask' - Type II error)
- **True Positives (TP): 492** (Images without masks correctly predicted as 'without mask')

Classification Report (Precision & Recall):

	precision	recall	f1-score	support
with_mask	0.87	0.98	0.92	560
without_mask	0.98	0.86	0.91	575
accuracy			0.92	1135
macro avg	0.93	0.92	0.92	1135
weighted avg	0.93	0.92	0.92	1135

- **'with_mask' class (0):**

- **Precision (0.87):** When the model predicts 'with mask', it is correct 87% of the time. This means there are some instances where it wrongly identifies a 'without mask' as 'with mask' (our False Negatives of 83).
- **Recall (0.98):** The model correctly identifies 98% of all actual 'with mask' cases. This indicates that the model is very good at not missing people who are wearing masks.

- **'without_mask' class (1):**

- **Precision (0.98):** When the model predicts 'without mask', it is correct 98% of the time. This means it is very reliable when it says someone is not wearing a mask.
- **Recall (0.86):** The model correctly identifies 86% of all actual 'without mask' cases. This means it misses some instances where people are not wearing masks, mistakenly classifying them as 'with mask' (our False Positives of 9).

Overall:

The model exhibits a strong overall accuracy of nearly 92%. It shows a higher recall for 'with_mask' (meaning it's good at detecting masks) and a higher precision for 'without_mask' (meaning its predictions of 'without mask' are highly trustworthy). The primary area for improvement seems to be reducing False Negatives for the 'without_mask' class, where people without masks are sometimes missed.

✓ 9. Make Predictions on New Images

```
import numpy as np
from tensorflow.keras.preprocessing import image
from google.colab import files
import matplotlib.pyplot as plt
import io
from PIL import Image # Import the Image module from PIL

# Function to preprocess the image for prediction
def preprocess_image(img_path):
    img = image.load_img(img_path, target_size=(128, 128))
    img_array = image.img_to_array(img)
    img_array = np.expand_dims(img_array, axis=0) # Add batch dimension
    img_array /= 255.0 # Normalize pixel values
```

```

    return img_array

# Upload an image
uploaded = files.upload()

for fn in uploaded.keys():
    # Display the uploaded image using PIL.Image.open
    img = Image.open(io.BytesIO(uploaded[fn]))
    plt.imshow(img)
    plt.title('Uploaded Image')
    plt.axis('off')
    plt.show()

    # Preprocess the uploaded image
    img_array = preprocess_image(io.BytesIO(uploaded[fn]))

    # Make prediction
    prediction = model.predict(img_array)

    # Interpret the prediction
    if prediction[0][0] > 0.5:
        result = "No Mask"
        confidence = prediction[0][0]
    else:
        result = "Mask"
        confidence = 1 - prediction[0][0]

    print(f"\nPrediction: {result} (Confidence: {confidence:.2f})")

```

masked person.jpg

masked person.jpg(image/jpeg) - 5155 bytes, last modified: 4/19/2026 - 100% done
 Saving masked person.jpg to masked person.jpg

Uploaded Image



1/1  0s 39ms/step

Prediction: Mask (Confidence: 0.71)

```

import numpy as np
from tensorflow.keras.preprocessing import image
from google.colab import files
import matplotlib.pyplot as plt
import io
from PIL import Image # Import the Image module from PIL

# Function to preprocess the image for prediction
def preprocess_image(img_path):
    img = image.load_img(img_path, target_size=(128, 128))
    img_array = image.img_to_array(img)
    img_array = np.expand_dims(img_array, axis=0) # Add batch dimension
    img_array /= 255.0 # Normalize pixel values

```



```
img_array = preprocess_image(img_array)
return img_array

# Upload an image
uploaded = files.upload()

for fn in uploaded.keys():
    # Display the uploaded image using PIL.Image.open
    img = Image.open(io.BytesIO(uploaded[fn]))
    plt.imshow(img)
    plt.title('Uploaded Image')
    plt.axis('off')
    plt.show()

    # Preprocess the uploaded image
    img_array = preprocess_image(io.BytesIO(uploaded[fn]))

    # Make prediction
    prediction = model.predict(img_array)

    # Interpret the prediction
    if prediction[0][0] > 0.5:
        result = "No Mask"
        confidence = prediction[0][0]
    else:
        result = "Mask"
        confidence = 1 - prediction[0][0]

    print(f"\nPrediction: {result} (Confidence: {confidence:.2f})")
```

non-masked person.jpg

non-masked person.jpg(image/jpeg) - 5017 bytes, last modified: 4/19/2026 - 100% done

Saving non-masked person.jpg to non-masked person.jpg

Uploaded Image



1/1 ————— 0s 33ms/step

Prediction: No Mask (Confidence: 0.98)

Start coding or [generate](#) with AI.

